

Target Specification, Not Model Capability, Explains Most Unit-Test Generation Failures at Medium Complexity

Nguyễn Tùng Ân, Đoàn Thị Thu Kim, Sơn Hải, Trương Hoàng Phúc, and Trần
Xuân Lộc

FPT University, Vietnam

Abstract. Keywords: Unit test generation · Large language models ·
Mutation testing · Cyclomatic complexity · Measurement validity

1 Introduction

A large language model asked to write a unit test for a specific function will often produce something that does not compile. It is tempting to read that as a limit of the model. In an earlier study of `gpt-4o-mini` on 120 medium-complexity functions we read it that way ourselves. This paper reports what we found when we examined the failures individually instead, and the answer turned out to be mostly about the benchmark and the measurements rather than about the model.

Two problems surfaced, and they compound.

The benchmark contains targets nothing can test. Sampling functions by cyclomatic complexity is standard, and it says nothing about whether a sampled function is reachable from an external test. In the ten projects studied here, 1,848 functions fall in the band $5 \leq CC \leq 10$; only 474 of them — 26% — are simultaneously public and unambiguously named. A generator that fails on the other 74% has been asked to do something no tool can do, and the resulting failure is recorded as a failure of the generator. Attributing a method to the wrong declaring class has the same effect and is easy to do silently: a brace counter that does not pop on block exit assigns a method to whichever nested type was declared most recently.

The usual metrics can be satisfied without testing anything. Compilation success and coverage are the two most commonly reported outcomes in this literature. Both can be obtained by a suite that verifies nothing. Reflection turns every identifier into a string, so a call naming a method that does not exist compiles with return code 0; and a generated suite in our own corpus consists of a single assertion-free call to the wrong function, which passes. The consequence is not that these metrics are noisy — it is that a number can be reported when the underlying measurement has stopped working, with nothing to signal it. We

encountered this seven times during the study, and four of those were defects we introduced ourselves while repairing the other three (Section 6).

This paper

We rebuild the benchmark under five testability criteria fixed in advance, sample 60 Java and 60 Python functions from the surviving pool, and measure every suite at four nested tiers, resting all claims on mutation score inside the target’s line range — the only tier that cannot be satisfied vacuously. Two *one-shot* prompt conditions differing in exactly one variable isolate the effect of supplying the target specification, and three established generators (EvoSuite, Randoop, Pynguin) provide reference points. The sampling criteria, tool versions, budgets and hypotheses were committed to a public repository before any measurement was run.

Three findings stand out.

1. **Target specification improves reach, not fault detection.** On Python, adding the fully qualified target, its signature, and an executed constructor call raises the proportion of suites that reach the target from 29/60 to 42/60 ($p = 0.003$, rank-biserial +0.61). The proportion that detects a fault rises from 16/60 to 22/60, which is *not* significant ($p = 0.33$). The prompt gets the test to the right place; it does not make the test check anything. A study reporting only coverage would have called this an unambiguous improvement.
2. **The LLM and the search-based tools win in different languages.** On Python, both prompt conditions beat Pynguin decisively — 2/60 for Pynguin against 16/60 and 22/60 — at $p \leq 0.003$ in every pairing, and this holds under both of the two Pynguin configurations we ran. On Java the ordering reverses, with both mature search-based tools ahead of the model. Neither result generalises to “LLMs are better” or “LLMs are worse.”
3. **Most of the attrition is structural.** Across both languages the dominant obstacle is not writing assertions but obtaining an instance to assert about. This is the same barrier the search-based literature identified for object-oriented code, and the model does not escape it.

We report the sampling funnel in full, both Pynguin configurations side by side, and both settings of a green-test gate whose Java implementation we found to disagree with its own pre-registered definition. Where a repair moved results in our own favour we say so explicitly and give the numbers under both settings, so that no conclusion in this paper depends on which of the two a reader prefers.

2 Related Work

2.1 Search-Based and Random Test Generation

Automated unit-test generation predates language models by more than a decade. Randoop builds test sequences incrementally and uses feedback from executing

each prefix to discard sequences that are illegal or redundant [11]. EvoSuite treats an entire test suite as one individual in a genetic search and evolves it against a coverage-based fitness function [3]. Both have been evaluated at scale: EvoSuite reached high branch coverage across a large sample of open-source classes, while also showing that a substantial fraction of classes remain hard for reasons unrelated to the search itself [4]. Pynguin ports the search-based approach to Python, where the absence of static types removes much of the information the Java generators depend on [8,9].

Two findings from this line of work bear directly on our design. First, coverage and fault detection are not interchangeable: Shamsiri et al. found that automatically generated suites with high coverage detected only a minority of real faults [14], and Inozemtseva and Holmes report that coverage correlates with suite effectiveness far more weakly than is commonly assumed [5]. Second, the difficulty of a target is not captured by its complexity alone — constructing the receiver object is frequently the binding constraint. We return to both points in Sections 3.3 and ??.

2.2 LLM-Based Test Generation

Code-trained language models made test generation a prompting problem [2], and the earliest work in this direction fine-tuned transformers on method–test pairs with surrounding focal context [16]. Subsequent systems moved to prompting. TestPilot generates tests for JavaScript functions from documentation and usage examples, and re-prompts with the failure message when a test does not pass [13]. ChatTester adds an iterative repair loop driven by compiler and runtime errors, reporting large gains over a single-shot baseline [18]. Siddiq et al. evaluate several models on JUnit generation and find compilation failure — not weak assertions — to be the dominant failure mode [15]. A recent survey catalogues the space and notes that evaluation practice is uneven across the field [17].

Our study is deliberately narrower than these systems in one respect and stricter in another. We use a *single* API call per function, with no repair loop and no retrieval, because a repair loop conflates the quality of the initial generation with the quality of the error signal fed back into it; when compilation failure is itself the object of study, that conflation is fatal. Where those systems ask how high the success rate can be driven, we ask what the initial generation is actually missing, and whether it is something a prompt can supply. The two prompt conditions of Section 3.2 differ in exactly one variable so that the answer is attributable to that variable.

2.3 Evaluation Metrics and Their Validity

Mutation analysis is the standard fault-based criterion for test-suite adequacy, and mutants have been shown to be a reasonable — if imperfect — substitute for real faults in controlled comparisons [6,12]. It is comparatively expensive, which is why much of the LLM test-generation literature reports compilation success and line or branch coverage instead.

This substitution is where our contribution sits. Compilation success and coverage are cheap because they ask less, and what they ask can be satisfied without testing anything: a reflective call compiles regardless of whether the named method exists, and a suite that constructs no object and calls nothing still passes. Neither weakness is hypothetical; both occur in the corpus measured here (Section 3.3). We therefore report four nested tiers and rest every claim on the only tier that cannot be satisfied vacuously.

A second and less discussed threat concerns the benchmark rather than the metric. Sampling targets by cyclomatic complexity [10] is standard practice, but complexity does not imply reachability: a method may be private, may be one of several overloads sharing a name, or may belong to a class that cannot be instantiated from outside. A suite that fails on such a target tells us about the sampling, not about the generator. To our knowledge, the fraction of a complexity-sampled benchmark that is actually addressable from an external test is rarely reported. We report the full sampling funnel for this reason, and the attrition turns out to be the study’s most transferable number.

2.4 Methodological Practice

Because the present work re-examines a dataset after having seen results obtained on an earlier one, the distinction between repairing an instrument and selecting a favourable outcome is central. Kerr’s account of hypothesising after the results are known [7] describes the failure mode precisely. We address it in two ways: the sampling criteria, tool versions, budgets and hypotheses were committed to the repository before any measurement was run, and the earlier study’s dataset and reported numbers are left unchanged — this paper stands beside it rather than replacing it. For the statistical comparisons we follow the guidance of Arcuri and Briand on non-parametric testing and effect sizes for randomised algorithms in software engineering [1].

3 Methodology

3.1 Dataset Construction

We mine functions from ten open-source repositories pinned at fixed commits (eight Java projects: `commons-cli`, `commons-csv`, `commons-collections`, `commons-math`, `gson`, `jsoup`, `joda-time`, `jfreechart`; two Python projects: `flask`, `requests`). Cyclomatic complexity is measured with Lizard 1.23.0 and we keep the medium band $5 \leq CC \leq 10$, matching the complexity range studied in our earlier work.

Complexity alone, however, does not make a function a valid test target. A study that samples on complexity but ignores *testability* measures two things at once: how well a generator writes tests, and how often the sampled target could have been tested at all. We therefore apply five additional criteria, each fixed before any test was generated and each justified by whether *any* tool — an LLM, a search-based generator, or a human — could reach the target through the public API.

Table 1. Sampling funnel. Each row reports how many candidates survive the criterion.

Stage	Criterion	Remaining
F0	$5 \leq CC \leq 10$ (Lizard)	1,848
F1	non-trivial body	1,848
F3	unambiguous name	753
F2	public / exported	474
F5	dynamically resolvable (Python)	64
	final pool	Java 406, Python 64
	sampled	Java 60, Python 60

F1 — Non-trivial body. $nloc \geq 3$. Removes stubs, abstract declarations and interface signatures, which contain no branches to cover.

F2 — Public or exported. Java: `public` only. Python: names not prefixed with an underscore. A `private` method cannot be invoked from an external test class by ordinary means; reflection can reach it, but then the compiler no longer checks any name in the test (Section 3.3), so compilation ceases to be evidence of anything.

F3 — Unambiguous name. Java: the method name occurs exactly once in its declaring class. Python: the function name occurs exactly once in its module. An overloaded name cannot be specified unambiguously in a prompt, and the compiler rejects the resulting call as ambiguous.

F4 — Verified declaring class. The enclosing class is resolved with a brace-depth stack that *pops* on block exit. Taking “the nearest enclosing class declaration” without popping attributes a method to whichever nested class was declared most recently, which is wrong whenever the target follows a nested type in the same file.

F5 — Dynamically resolvable (Python). The module is imported, the qualified name is walked with `getattr`, and the result is confirmed callable. Syntactic inference is not evidence that a target can be reached at run time.

Table 1 reports the full funnel. We report every stage rather than only the final count, because the attrition is itself a finding: of 1,848 functions in the complexity band, only 474 (26%) are simultaneously public and unambiguous. Roughly three quarters of medium-complexity functions in these projects are not valid targets for an external test generator.

From the surviving pool we sample 60 Java and 60 Python functions, balanced across repositories (Java: 3–9 functions per project; Python: 32 from `requests`, 28 from `flask`). Complexity is concentrated at the lower end of the band ($CC=5$: 56 functions; $CC \geq 8$: 23), and the median function is 17 lines. Of the Python targets, 34 are instance methods and 26 are module-level functions. Every one of the 120 targets was confirmed resolvable at run time before generation began.

3.2 Test Generation

All tests are produced by `gpt-4o-mini-2024-07-18` at `temperature = 0`, `top_p = 1.0`, `max_tokens = 2048`, in a **single API call per function** with no feedback loop and no retrieval.

We compare two *one-shot* prompt conditions. Both contain exactly one worked exemplar (a trivial `add` function with its test), so both are one-shot in the standard sense, and both issue one call. They differ in exactly one respect:

- V1 — baseline.** The prompt names the target function and shows its source. This is the prompt used in our earlier study.
- V2 — target-specified.** Identical to V1, with one additional block inserted before the task: the fully qualified target name, its exact signature, its declaring class, and — for instance methods — a constructor call that has been *executed* and confirmed to build the receiver successfully.

Holding the exemplar fixed matters. An earlier draft of the enriched prompt was rewritten from scratch and, in the process, dropped the exemplar; it therefore differed from V1 in two respects at once (no exemplar, plus target metadata) and no causal attribution was possible from it. We report that variant only as a secondary observation and label it zero-shot, since it is one.

3.3 Measurement: Four Tiers

A single “valid / invalid” flag conflates outcomes that differ in kind, so we report four nested tiers. Each answers a different question, and — importantly — the first two can be satisfied by a test that verifies nothing.

- T1 — Compiles / collects.** The test file is syntactically and type-correct.
- T2 — At least one green test.** At least one test passes against the unmodified source.
- T3 — Reaches the target.** Branch coverage > 0 within the target’s line range $[start, end]$.
- T4 — Detects faults.** Mutation score > 0 within the same line range.

Only T4 supports the claim that a suite *tests* the target. We show that T1 and T2 are both reachable without testing anything:

- **T1 under reflection.** Reflection turns every name into a string, so the compiler stops checking it. A call to `getDeclaredMethod("aNameThatDoesNotExist")` compiles with return code 0 and fails only at run time. Any compilation-based validity rate is therefore uninformative for suites that use reflection.
- **T2 under an assertion-free suite.** The search-based suite generated for `requests.hooks` (target `dispatch_hook`) contains exactly one test, and that test is a single bare call to a *different* function in the same module, with no assertion of any kind. It compiles, it collects, and it passes — T1 and T2 are both satisfied — while measured branch coverage inside the target’s line range is 0. A green-test criterion records this as a success.

The two tiers are therefore not merely weaker than T4; they can be satisfied by suites that are, respectively, unchecked by the compiler and empty of assertions.

Branch coverage is measured with `coverage.py` (Python) and by executing the generated JUnit suites through the JUnit Platform launcher (Java), in both cases attributed to the target’s line range rather than to the whole file. Mutation analysis applies the same operator set to both languages — arithmetic ($+/-$, $\times/\div$), relational ($> / \geq$, $< / \leq$), equality ($= / \neq$), boolean ($\wedge / \vee$), boolean constant inversion, and numeric constant increment — capped at 20 mutants per function, restricted to the target’s line range.

Two procedural safeguards are worth stating because their absence silently corrupts results. First, mutation requires overwriting the source; we mutate a *copy* of the source tree placed ahead of the original on the import path, so that concurrent measurements reading the same tree are unaffected. Second, tool versions must be matched across families: JUnit 5 pairs Jupiter 5.x with Platform 1.x, and mixing a Jupiter 6 artifact with JUnit 5-style tests makes the launcher discover zero tests while exiting successfully — a failure that reports as a legitimate zero.

3.4 Baselines

Each baseline generates per class (Java) or per module (Python), matching how the tools operate; results are then attributed to individual functions by line range, exactly as for the LLM suites.

EvoSuite 1.2.0 , 60 s per class, run under JDK 11. EvoSuite performs deep reflection into JDK internals, which the module system blocks from Java 9 onward, and its master process launches its client from `JAVA_HOME` — so both must point at the same JDK, or the client dies while the master still exits successfully. Because the projects were compiled with JDK 17, we recompiled sources to Java 11 bytecode into a separate output directory, leaving the original build untouched.

Randoop 4.3.3 , 60 s per class, JDK 17.

Pynguin 0.45.0 , 90 s per module, under **two configurations reported side by side**. In its default configuration Pynguin serialises module state to a worker process with `dill` and aborts on C-extension type objects that cannot be resolved by name, producing no output at all for several modules. We therefore also run it with the master–worker split disabled. Reporting both separates “the tool scored low” from “the tool never ran”.

3.5 Statistical Analysis

Conditions are compared with the paired Wilcoxon signed-rank test, matching on function identifier, at $\alpha = 0.05$, with matched-pairs rank-biserial correlation as effect size. Java and Python are analysed separately throughout: the toolchains differ, and Java enforces access control while Python does not, so their failure

modes are not commensurable. Every rate is reported over the 60 functions of its own language.

Results are given at two levels in parallel: *all* ($n = 60$ per language, with unmeasurable cases scored 0) and *effective* (the subset with a non-zero value). The *all* level is reported even where the *effective* level is more favourable.

All configuration decisions above — criteria, tool versions, budgets, the two-branch Pynguin design, and the hypotheses — were recorded in a pre-registration committed before any baseline was executed.

4 results

5 discussion

6 Threats to Validity

6.1 Construct Validity

What a validity flag measures. The central construct threat in this area is that the usual “valid suite” indicator does not measure what its name suggests, and fails silently when it does not. We encountered four distinct instances, three of them in our own instrumentation.

First, a compilation flag inferred from a coverage report measures only whether the class appears in that report. Coverage tools enumerate every class in a project, including classes that were never loaded; such a class yields zero covered branches, which the parser then records as a successfully compiled function with zero coverage. Second, a criterion of “at least one test executed” counts a file in which every test fails. Third, reflection removes the compiler’s ability to check names, so a compilation-based rate is uninformative for any suite that uses it. Fourth, mismatched tool versions can make a test launcher discover zero tests while exiting successfully.

None of these produce an error. Each returns a plausible number. We therefore report the four nested tiers of Section 3.3 rather than any single flag, and treat only T4 — fault detection within the target’s line range — as evidence that a suite tests its target.

Mutation operators. Our operator set is deliberately small and identical across both languages, so that the two are measured with the same instrument. It is weaker than a production mutation framework, so absolute mutation scores are not comparable with studies using PIT or `mutmut`; only within-study comparisons are meaningful. We also cap mutants at 20 per function, which bounds the resolution of the score for larger functions.

6.2 Internal Validity

Sampling. Our own miner mis-attributed one Java target: `LocalTime.getField` is declared `protected`, but the parser matched a same-named declaration in a

nested class and recorded it as public, so it passed the public-only filter. An audit against the real declaration line found this to be the only such case in 60 Java functions; we exclude it from the analysis. We record it here because it is precisely the class-attribution error this study is about, reproduced in our own tooling — knowing about a failure mode does not confer immunity to it.

Concurrency. Mutation analysis overwrites source files. Running it directly on a shared source tree corrupts any concurrent measurement reading that tree, silently. We mutate an isolated copy placed ahead of the original on the import path, and verify after each run that the original tree is unmodified.

6.3 External Validity

Source exhaustion in Python. After all testability criteria, only 64 Python functions in the two projects qualify, of which 34 already appeared in our earlier dataset. A strictly held-out Python replication is therefore limited to roughly 30 functions without adding new repositories. Our Python sample is also drawn from two web-oriented libraries and may not generalise to other domains.

Complexity distribution. Although the band is $5 \leq CC \leq 10$, the sample concentrates at the lower end (56 of 120 functions have $CC = 5$). Conclusions about the upper half of the band rest on fewer observations.

Model and configuration. A single model at a single temperature was used. We do not claim the findings transfer to other models, and the deterministic setting means we do not estimate run-to-run variance.

6.4 Conclusion Validity

Baseline configuration. Three limitations constrain how far the baseline comparison can be pushed. EvoSuite was run with a classpath containing only the module’s own compiled classes: with the full dependency classpath its bytecode reader aborts in a way that produces no output while reporting success, and we were unable to identify the specific offending artifact. EvoSuite may therefore be weaker here than in a fully configured environment. Seven Java targets in one multi-module project could not be recompiled to Java 11 bytecode and fall outside EvoSuite’s coverage. Finally, our Randoop version differs from the one used in our earlier study, so Randoop numbers are compared only within this study.

Pre-registration and one deviation. Research questions, hypotheses, tool versions, budgets and the two-branch Pynguin design were committed before any baseline was run. One condition was added afterwards: on inspection, the enriched prompt we had planned to use turned out to lack the worked exemplar, making it zero-shot rather than one-shot and differing from the baseline in two respects at once. We added a corrected condition that holds the exemplar fixed and

varies only the target specification, recorded the change in the pre-registration with its date, and left the original hypothesis unaltered. The uncorrected variant is reported only as a secondary observation.

7 conclusion

References

1. Arcuri, A., Briand, L.: A hitchhiker’s guide to statistical tests for assessing randomized algorithms in software engineering. *Software Testing, Verification and Reliability* **24**(3) (2014)
2. Chen, M., Tworek, J., Jun, H., Yuan, Q., et al.: Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374* (2021)
3. Fraser, G., Arcuri, A.: EvoSuite: Automatic test suite generation for object-oriented software. In: *Proceedings of the 19th ACM SIGSOFT Symposium and the 13th European Conference on Foundations of Software Engineering (ESEC/FSE)* (2011)
4. Fraser, G., Arcuri, A.: A large-scale evaluation of automated unit test generation using EvoSuite. *ACM Transactions on Software Engineering and Methodology (TOSEM)* **24**(2) (2014)
5. Inozemtseva, L., Holmes, R.: Coverage is not strongly correlated with test suite effectiveness. In: *Proceedings of the 36th International Conference on Software Engineering (ICSE)* (2014)
6. Just, R., Jalali, D., Inozemtseva, L., Ernst, M.D., Holmes, R., Fraser, G.: Are mutants a valid substitute for real faults in software testing? In: *Proceedings of the 22nd ACM SIGSOFT International Symposium on Foundations of Software Engineering (FSE)* (2014)
7. Kerr, N.L.: HARKing: Hypothesizing after the results are known. *Personality and Social Psychology Review* **2**(3) (1998)
8. Lukasczyk, S., Fraser, G.: Pynguin: Automated unit test generation for Python. In: *Proceedings of the 44th International Conference on Software Engineering: Companion Proceedings (ICSE Companion)* (2022)
9. Lukasczyk, S., Kroiß, F., Fraser, G.: An empirical study of automated unit test generation for Python. *Empirical Software Engineering* **28**(2) (2023)
10. McCabe, T.J.: A complexity measure. *IEEE Transactions on Software Engineering* **SE-2**(4) (1976)
11. Pacheco, C., Lahiri, S.K., Ernst, M.D., Ball, T.: Feedback-directed random test generation. In: *Proceedings of the 29th International Conference on Software Engineering (ICSE)* (2007)
12. Papadakis, M., Kintis, M., Zhang, J., Jia, Y., Le Traon, Y., Harman, M.: Mutation testing advances: An analysis and survey. *Advances in Computers* **112** (2019)
13. Schäfer, M., Nadi, S., Eghbali, A., Tip, F.: An empirical evaluation of using large language models for automated unit test generation. *IEEE Transactions on Software Engineering* **50**(1) (2024)
14. Shamshiri, S., Just, R., Rojas, J.M., Fraser, G., McMinn, P., Arcuri, A.: Do automatically generated unit tests find real faults? an empirical study of effectiveness and challenges. In: *Proceedings of the 30th IEEE/ACM International Conference on Automated Software Engineering (ASE)* (2015)

15. Siddiq, M.L., Santos, J.C.S., Tanvir, R.H., Ulfat, N., Al Rifat, F., Carvalho Lopes, V.: Using large language models to generate JUnit tests: An empirical study. In: Proceedings of the 28th International Conference on Evaluation and Assessment in Software Engineering (EASE) (2024)
16. Tufano, M., Drain, D., Svyatkovskiy, A., Deng, S.K., Sundaresan, N.: Unit test case generation with transformers and focal context. arXiv preprint arXiv:2009.05617 (2020)
17. Wang, J., Huang, Y., Chen, C., Liu, Z., Wang, S., Wang, Q.: Software testing with large language models: Survey, landscape, and vision. *IEEE Transactions on Software Engineering* **50**(4) (2024)
18. Yuan, Z., Lou, Y., Liu, M., Ding, S., Wang, K., Chen, Y., Peng, X.: No more manual tests? evaluating and improving ChatGPT for unit test generation. arXiv preprint arXiv:2305.04207 (2023)